

## Gradient Descent Algorithm for Linear Regression (10 points)

The gradient descent algorithm is a powerful optimization technique used to iteratively refine model parameters for a better fit. In the context of linear regression, it aims to find optimal values of the slope  $m$  and intercept  $c$  that minimize the mean squared error between the predicted and actual values.

The steps for each iteration (or epoch) are:

1. **Prediction:** Use the current values of  $m$  and  $c$  to calculate the predicted values:

$$y_{\text{pred}} = w \cdot x + b$$

2. **Compute Gradients:** Determine the gradient of the loss with respect to each parameter:

- Gradient with respect to the slope  $m$ :

$$D_w = \frac{-2}{n} \sum (x \cdot (y_{\text{noisy}} - y_{\text{pred}}))$$

- Gradient with respect to the intercept  $c$ :

$$D_b = \frac{-2}{n} \sum (y_{\text{noisy}} - y_{\text{pred}})$$

3. **Update Parameters:** Adjust  $w$  and  $b$  based on the gradients and the learning rate  $\alpha$ :

$$w = w - \alpha \cdot D_w$$

$$b = b - \alpha \cdot D_b$$

In the code below replace `$# Your code here$` with your code

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import random
import os
import imageio

random.seed(42)
np.random.seed(42)

x = np.linspace(0, 10, 100)

# Noisy Linear Data generation
y_linear_noisy = 3*x + 2 + np.random.randn(100)*5
```

```

# Initial values for m (slope) and c (intercept)
w, b = 0, 0 # Model parameters
learning_rate = 0.01 # learning rate
epochs = 15 # Limit to 15 epochs

# Lists to save the plots at each epoch for visualization
lines = []

# print variables to observe
print(f"x: {x}")
print()
print(f"y_linear_noisy: {y_linear_noisy}")

x: [ 0.          0.1010101  0.2020202  0.3030303  0.4040404
0.50505051
 0.60606061  0.70707071  0.80808081  0.90909091  1.01010101
1.11111111
 1.21212121  1.31313131  1.41414141  1.51515152  1.61616162
1.71717172
 1.81818182  1.91919192  2.02020202  2.12121212  2.22222222
2.32323232
 2.42424242  2.52525253  2.62626263  2.72727273  2.82828283
2.92929293
 3.03030303  3.13131313  3.23232323  3.33333333  3.43434343
3.53535354
 3.63636364  3.73737374  3.83838384  3.93939394  4.04040404
4.14141414
 4.24242424  4.34343434  4.44444444  4.54545455  4.64646465
4.74747475
 4.84848485  4.94949495  5.05050505  5.15151515  5.25252525
5.35353535
 5.45454545  5.55555556  5.65656566  5.75757576  5.85858586
5.95959596
 6.06060606  6.16161616  6.26262626  6.36363636  6.46464646
6.56565657
 6.66666667  6.76767677  6.86868687  6.96969697  7.07070707
7.17171717
 7.27272727  7.37373737  7.47474747  7.57575758  7.67676768
7.77777778
 7.87878788  7.97979798  8.08080808  8.18181818  8.28282828
8.38383838
 8.48484848  8.58585859  8.68686869  8.78787879  8.88888889
8.98989899
 9.09090909  9.19191919  9.29292929  9.39393939  9.49494949
9.5959596
 9.6969697  9.7979798  9.8989899  10.          ]

y_linear_noisy: [ 4.48357077  1.6117088  5.8445033  10.52424019
2.04135434  2.34446673

```

11.7142459	7.95838577	2.07687049	7.44007295	2.71321457
3.00468457				
6.84617499	-3.62700728	-2.38216492	3.7340169	1.78432925
8.72275181				
2.91442508	0.69605725	15.38884991	7.23475486	9.00430769
1.84595604				
6.55081365	10.13037052	4.12381999	12.06030827	7.48165504
9.32941004				
8.08237603	20.65533032	11.62948357	6.71144536	16.41575486
6.50184236				
13.95340888	3.41377059	6.87422127	14.802488	17.81354502
15.28108383				
14.14903132	13.52478455	7.94072338	12.03714259	13.63620008
21.52803537				
18.26354599	8.03328407	18.771935	15.52913405	14.37296576
21.1189875				
23.51863398	23.32306726	14.77360935	17.72666539	21.23207473
24.75651351				
17.78594699	19.5565536	15.25620392	15.10987597	25.45656851
28.47816984				
21.63994939	27.32069479	24.41424073	19.68349214	25.01909924
31.20533435				
23.63905162	31.9444304	11.3255169	28.83678525	25.46553837
23.83829658				
26.09516752	16.00154937	25.1440648	28.3310174	34.23795507
24.56016406				
23.41207744	25.24879054	32.63761665	30.00739191	26.01786565
31.53603414				
29.75811502	34.41898253	26.36852241	28.54350745	28.52430772
23.47030405				
32.57151048	32.69921575	31.72253698	30.82706433	

```

0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0.
0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0.
0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.
0.
0. 0. 0. 0.]
-216.99043507389993 -32.961534826059044
2.1699043507389995 0.32961534826059047

```

```

# Gradient Descent

```

```

for epoch in range(epochs):

```

```

    # Make predictions, watch formula listed above to implement

```

```

    y_pred = w * x + b

```

```

    # Calculate gradients

```

```

    D_w = (-2/len(x)) * sum(x * (y_linear_noisy - y_pred))

```

```

    D_b = (-2/len(x)) * sum(y_linear_noisy - y_pred)

```

```

    # Update parameters, watch formulas listed above to implement

```

```

    w = w - learning_rate * D_w

```

```

    b = b - learning_rate * D_b

```

```

    # Visualization

```

```

    plt.figure(figsize=(10,5))

```

```

    plt.scatter(x, y_linear_noisy, c='blue', label='Noisy Linear
Data')

```

```

    plt.plot(x, w*x+b, '-r', label=f'Epoch {epoch+1}')

```

```

    plt.xlabel('X')

```

```

    plt.ylabel('Y')

```

```

    plt.title(f'Gradient Descent Epoch {epoch+1}')

```

```

    plt.legend()

```

```

    plt.grid(True)

```

```

    filename = f'epoch_{epoch}.png'

```

```

    plt.savefig(filename)

```

```

    lines.append(filename)

```

```

    plt.close()

```

```

# Create gif with slowed down duration and infinite loop

```

```

with imageio.get_writer('gradient_descent.gif', mode='I',

```

```

duration=700, loop = 0) as writer: # duration is in seconds per frame

```

```

    for filename in lines:

```

```

        image = imageio.imread(filename)

```

```

        writer.append_data(image)

```

```

# Cleanup the individual frames

```

```

for filename in lines:

```

```

    os.remove(filename)

```

```
# Display the optimized parameters and the gif
print(f"Optimized Parameters: Slope (m) = {w}, Intercept (c) = {b}")

# Display gif in Jupyter notebook
# from IPython.display import Image
# Image(filename="gradient_descent.gif") # loop=0 makes it loop indefinitely

/var/folders/b8/bxlpngnn4dd74v10s4rp5mp40000gn/T/
ipykernel_39157/918653717.py:31: DeprecationWarning: Starting with
ImageIO v3 the behavior of this function will switch to that of
iio.v3.imread. To keep the current behavior (and make this warning
disappear) use `import imageio.v2 as imageio` or call
`imageio.v2.imread` directly.
  image = imageio.imread(filename)

Optimized Parameters: Slope (m) = 3.1605497473212765, Intercept (c) =
0.526842024174439
```

## Linear regression using sklearn (10 points)

Load the dataset and perform basic data exploration.

```
from sklearn.datasets import fetch_california_housing
import pandas as pd

california_housing = fetch_california_housing(as_frame=True)

# Load dataset
df = pd.DataFrame(california_housing.data,
                  columns=california_housing.feature_names)
df['MedHouseVal'] = california_housing.target

# Display the first few rows
print(df.head())
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467

Longitude    MedHouseVal

0	-122.23	4.526
1	-122.22	3.585
2	-122.24	3.521
3	-122.25	3.413
4	-122.25	3.422

Split the data into training and testing sets.

```
from sklearn.model_selection import train_test_split

X = df.drop('MedHouseVal', axis=1)
y = df['MedHouseVal']

print(f"x: {type(x)}")
print(f"y: {type(y)}")

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

print(f"X_train: {type(X_train)}")
print(f"Y_train: {type(y_train)}")

x: <class 'numpy.ndarray'>
y: <class 'pandas.core.series.Series'>
X_train: <class 'pandas.core.frame.DataFrame'>
Y_train: <class 'pandas.core.series.Series'>
```

Fit a linear regression model to the training data and evaluate its performance on the testing set.

```
from sklearn.linear_model import SGDRegressor
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
import random

# set random seed
random.seed(42)
np.random.seed(42)

# Create a model object
reg = SGDRegressor(max_iter=40) #Your code here (define linear
regression model with max_iter parameter equal to 40)
model = make_pipeline(StandardScaler(),reg)
# In case of any difficulties check:
https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDRegressor.html#sklearn.linear\_model.SGDRegressor

# Fit the model to the training data
# Your code here
```

```

model.fit(X_train, y_train)


# Make predictions on the test set
y_pred = model.predict(X_test) # Your code here

# Calculate and display the mean squared error between the actual and
# predicted values on test set
mse = mean_squared_error(y_test, y_pred)
# acc = accuracy_score(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
# print(f"Accuracy: {acc}")

# Plot the actual vs predicted values
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Prices vs Predicted Prices")
plt.show()

Mean Squared Error: 0.550598777585777

```



Tip: SGDRegressor has fit and predict methods

(10 points)

## Changing hyperparameters

What happens with accuracy if you change `max_iter` to 3? Modify and place code from the previous cell in the cell below.

Tip: Examples of how to use `max_iter` can be found here:

[https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.SGDRegressor.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.SGDRegressor.html)

How can you explain this behavior? (5 points):

```
# Place modified code here

# Create a model object
reg = SGDRegressor(max_iter=3) #Your code here (define linear
regression model with max_iter parameter equal to 40)
model = make_pipeline(StandardScaler(), reg)
# In case of any difficulties check:
https://scikit-learn.org/stable/modules/generated/sklearn.linear_model
.SGDRegressor.html#sklearn.linear_model.SGDRegressor

# Fit the model to the training data
# Your code here

model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = model.predict(X_test) # Your code here


# Calculate and display the mean squared error between the actual and
predicted values on test set
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")

# Plot the actual vs predicted values
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual Prices vs Predicted Prices")
plt.show()

Mean Squared Error: 1.2322400452788693

/Users/liuyucheng/anaconda3/envs/py38/lib/python3.8/site-packages/
sklearn/linear_model/_stochastic_gradient.py:1561: ConvergenceWarning:
Maximum number of iteration reached before convergence. Consider
increasing max_iter to improve the fit.
  warnings.warn(
```





Answer:

By changing the max\_iter to 3, I find that the slope of the point trend become lower, which means that the accuracy of the model become lower. As far as I am concerned, the model trained with 3 epoch is still in underfitting. Hence, the accuracy is still low in this case.

## Regularization Parameter Search for Ridge Regression (5 points)

Manually search for the best regularization parameter alpha in Ridge regression.

The goal of **Ridge regression** (aka Tikhonov regression) is to find a linear model that not only fits the data well but also has small parameter values. This yields to simpler models and therefore prevents overfitting to the data

The regularization parameter  $\alpha$  controls the trade-off: When  $\alpha = 0$ , Ridge regression becomes equivalent to linear regression without any regularization. As  $\alpha$  increases, the impact of the regularization term also increases, leading to a stronger preference for smaller parameter values (thus simpler models). If  $\alpha$  is very large, the regularization term dominates, causing the model parameters to tend towards zero, which leads to a model that is more biased. It's important to choose  $\alpha$  properly to balance a balance between underfitting and overfitting.

First step is Feature scaling

Feature scaling is a technique used in machine learning and statistics to normalize the range of independent variables or features of the data. This normalization often helps in speeding up the training process and can lead to better performance in certain algorithms that are sensitive to the scale of features.

StandardScaler transforms the data into a distribution with a mean of 0 and a standard deviation of 1.

```
from sklearn.preprocessing import StandardScaler

# Initialize and fit the scaler
scaler = StandardScaler()

# Fit the scaler to the training data and transform it to have zero mean and unit variance
X_train_scaled = scaler.fit_transform(X_train)

# Use the scaler fitted on the training data to scale the test data
X_test_scaled = scaler.transform(X_test)
```

Find optimal value of alpha for Ridge Regression

```
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error

# List of alphas (regularization parameter values) to test
alphas = [1e-5, 1e-3, 1e-2, 100, 500, 1000, 10000]

alpha = alphas[-1] # Your code here, you can chose any value from the list above

ridge = Ridge(alpha=alpha)
ridge.fit(X_train_scaled, y_train)
# In case of any difficulties check:
https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.Ridge.html#sklearn.linear\_model.Ridge

# Predict on the validation data
y_pred = ridge.predict(X_test) # Your code here

# Calculate the mean squared error
mse = mean_squared_error(y_test, y_pred)

print(f"Alpha Value: {alpha}")
print(f"Mean Squared Error: {mse}")

Alpha Value: 10000
Mean Squared Error: 113.45896636929982
```

```
/Users/liuyucheng/anaconda3/envs/py38/lib/python3.8/site-packages/sklearn/base.py:458: UserWarning: X has feature names, but Ridge was fitted without feature names
  warnings.warn(
```

y\_train

```
14196    1.030
8267     3.821
17445    1.726
14265    0.934
2271     0.965
```

```
...
```

```
11284    2.292
11964    0.978
5390     2.221
860      2.835
15795    3.250
```

```
Name: MedHouseVal, Length: 16512, dtype: float64
```

What happens to accuracy when alpha (regularization parameter) is too high or too low?

Your answer here

```
MSEs = []
for alpha in alphas:
    ridge = Ridge(alpha=alpha)
    ridge.fit(X_train_scaled, y_train)

    # Predict on the validation data
    y_pred = ridge.predict(X_test) # Your code here

    # Calculate the mean squared error
    mse = mean_squared_error(y_test, y_pred)
    MSEs.append(mse)
plt.plot(alphas, MSEs)
plt.show()
```

```
/Users/liuyucheng/anaconda3/envs/py38/lib/python3.8/site-packages/sklearn/base.py:458: UserWarning: X has feature names, but Ridge was fitted without feature names
```


```
  warnings.warn(
```

```
/Users/liuyucheng/anaconda3/envs/py38/lib/python3.8/site-packages/sklearn/base.py:458: UserWarning: X has feature names, but Ridge was fitted without feature names
```

```
  warnings.warn(
```

```
/Users/liuyucheng/anaconda3/envs/py38/lib/python3.8/site-packages/sklearn/base.py:458: UserWarning: X has feature names, but Ridge was
```

```
fitted without feature names
warnings.warn(
/Users/liuyucheng/anaconda3/envs/py38/lib/python3.8/site-packages/
sklearn/base.py:458: UserWarning: X has feature names, but Ridge was
fitted without feature names
warnings.warn(
/Users/liuyucheng/anaconda3/envs/py38/lib/python3.8/site-packages/
sklearn/base.py:458: UserWarning: X has feature names, but Ridge was
fitted without feature names
warnings.warn(
/Users/liuyucheng/anaconda3/envs/py38/lib/python3.8/site-packages/
sklearn/base.py:458: UserWarning: X has feature names, but Ridge was
fitted without feature names
warnings.warn(
/Users/liuyucheng/anaconda3/envs/py38/lib/python3.8/site-packages/
sklearn/base.py:458: UserWarning: X has feature names, but Ridge was
fitted without feature names
warnings.warn(
```



## Answer:

When the Alpha is too high, the model will try it best to minimize the parameters in the linear model. Hence, the parameter in the model may be minimized to zero. The accuracy is lower. On the other hand, when the Alpha is too low, the loss term of regularization will be extremely low, hence, it cannot constrain the parameters in the model.

# Classification with the Iris Dataset Using Logistic Regression (8 points)

Load the dataset and perform basic data exploration.

Split the data into training and testing sets.

Train a Logistic Regression classifier and evaluate its performance.

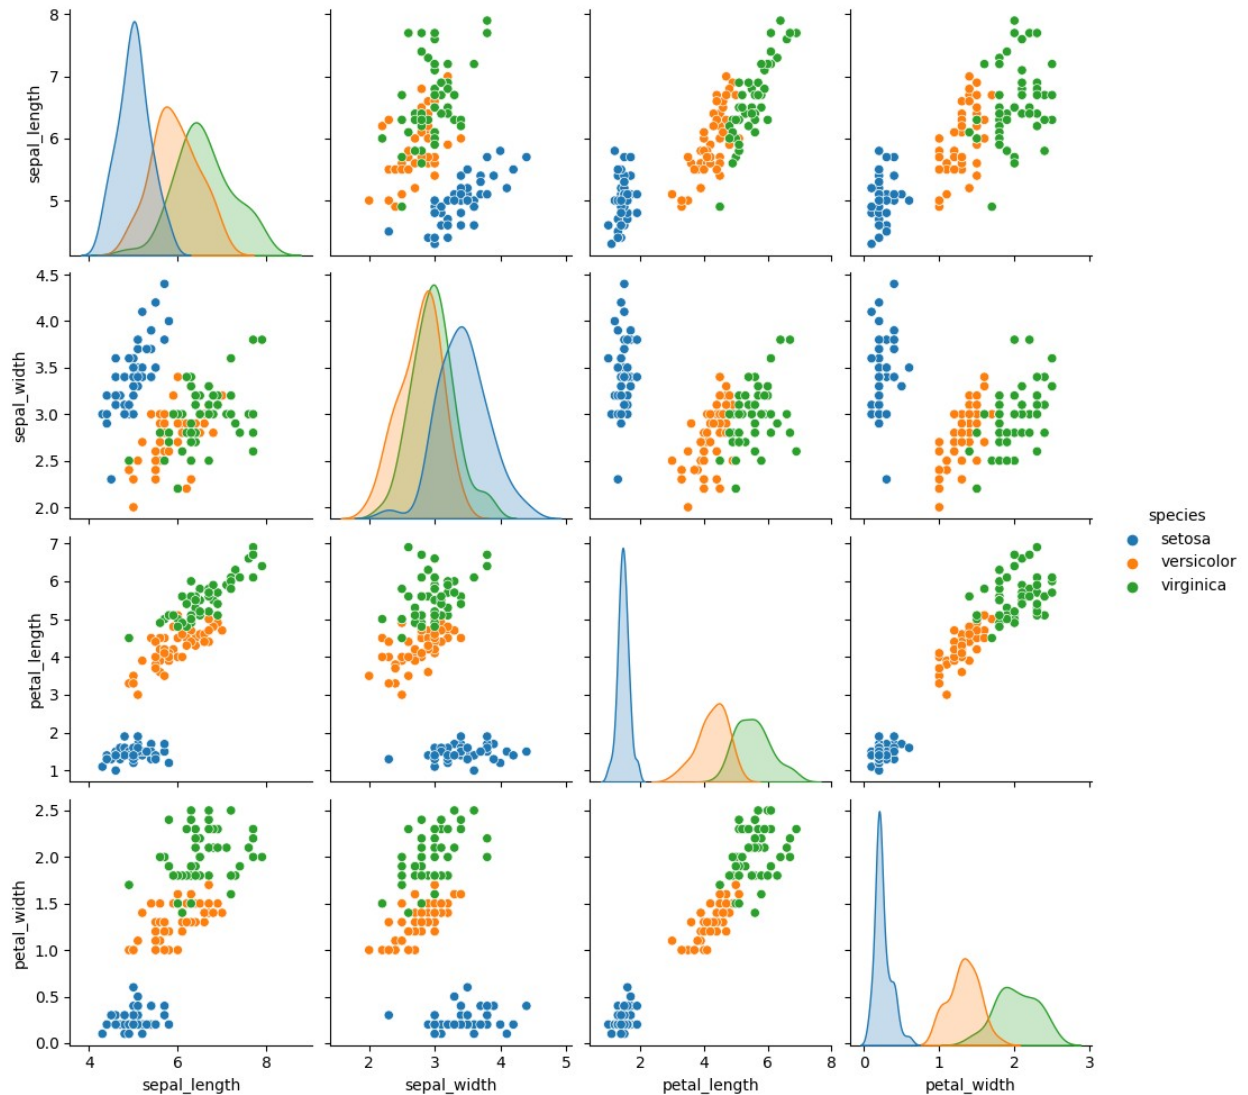
```
import seaborn as sns
import matplotlib.pyplot as plt

# Load the dataset
iris = sns.load_dataset("iris")

# Visualize the dataset using a pair plot
sns.pairplot(iris, hue="species")
plt.show()
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# Split the data into features and labels
X = iris.drop("species", axis=1)
y = iris["species"]

# Split the data into training and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)
```



```
# Standardize the features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

from sklearn.linear_model import LogisticRegression

# Create a logistic regression model with OvR multi-class strategy,
# peek suitable value for max_iter
clf = LogisticRegression(max_iter=100)# Your code here

# In case of any difficulties check:
https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.LogisticRegression.html#sklearn.linear\_model.LogisticRegression
# Train the model
# Your code here
```

```

clf.fit(X_train, y_train)

from sklearn.metrics import accuracy_score, classification_report

# Predict the species for the test set
y_pred = clf.predict(X_test) # Your code here

# Calculate the accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.4f}")

# Display a classification report
report = classification_report(y_test, y_pred)
print(report)

import numpy as np

# For simplicity, let's visualize using only the first two features
(sepal length and sepal width)
X_train_2d = X_train[:, :2]

# Train the model again on the 2D data
clf.fit(X_train_2d, y_train)

# Plot the decision boundaries
x_min, x_max = X_train_2d[:, 0].min() - 1, X_train_2d[:, 0].max() + 1
y_min, y_max = X_train_2d[:, 1].min() - 1, X_train_2d[:, 1].max() + 1

xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                     np.arange(y_min, y_max, 0.01))

Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])
Z = Z.reshape(xx.shape)

from sklearn.preprocessing import LabelEncoder

# Encode labels into numbers
label_encoder = LabelEncoder()
y_train_encoded = label_encoder.fit_transform(y_train)

# Train the classifier again on the 2D data
clf.fit(X_train_2d, y_train_encoded)

# Predict
Z = clf.predict(np.c_[xx.ravel(), yy.ravel()])

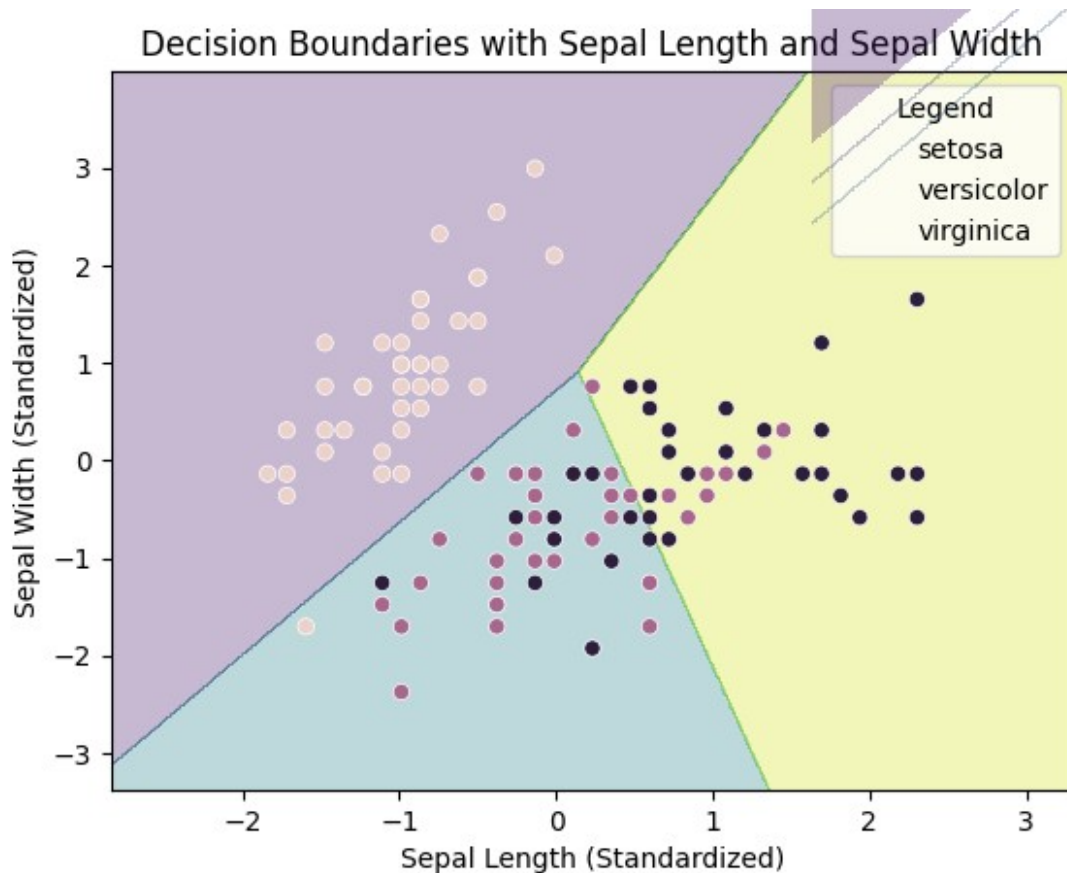
# No need to decode since we're just plotting the boundaries
Z = Z.reshape(xx.shape)

```

```
plt.contourf(xx, yy, Z, alpha=0.3)
sns.scatterplot(x=X_train_2d[:, 0], y=X_train_2d[:, 1],
hue=label_encoder.transform(y_train))
plt.title('Decision Boundaries with Sepal Length and Sepal Width')
plt.xlabel('Sepal Length (Standardized)')
plt.ylabel('Sepal Width (Standardized)', )
plt.legend(title='Legend', loc='best', labels=['setosa', 'versicolor',
'veirginica'])
plt.show()
```

Accuracy: 1.0000

	precision	recall	f1-score	support
setosa	1.00	1.00	1.00	10
versicolor	1.00	1.00	1.00	9
virginica	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30





# Understanding Overfitting and Underfitting (10 points)

One of the most important step in evaluating a machine learning model is to undersand whether the model overfits or underfits the data.

```
import numpy as np
from sklearn.pipeline import Pipeline
from sklearn.model_selection import cross_val_score
from sklearn import preprocessing
import matplotlib.pyplot as plt
from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression
# set seed
np.random.seed(132)

def true_fun(X):
    return np.cos(2 * np.pi * X)

n_samples = 50
X = np.sort(np.random.rand(n_samples))
y = true_fun(X) + np.random.randn(n_samples) * 0.2
```

The aim is to try to fit a polynomial function of degree  $n$  to the data.  $n$  should be chosen as to not overfit or underfit the underlying data. This is often a trial-and-error process, using cross-validation and visualization to estimate the "best" fit to the trainig data.

The first step is to fit the data.

```
degrees = int(input("Please Enter The Degree of Polynomial Between 1-20:"))

polynomial_features = PolynomialFeatures(degree=degrees,
include_bias=False)
linear_regression = LinearRegression()
pipeline = Pipeline([("polynomial_features", polynomial_features),
                    ("linear_regression", linear_regression)])
pipeline.fit(X[:, np.newaxis], y)


Pipeline(steps=[('polynomial_features',
PolynomialFeatures(degree=10, include_bias=False)),
('linear_regression', LinearRegression())])
```

Then to evaluate the model using cross-validation and visualization.

```
plt.figure(figsize=(15, 6))

scores = cross_val_score(pipeline, X[:, np.newaxis], y,
                          scoring="neg_mean_squared_error", cv=10)

X_test = np.linspace(0, 1, 100)
plt.plot(X_test, pipeline.predict(X_test[:, np.newaxis]),
         label="Model")
plt.plot(X_test, true_fun(X_test), label="True function")
plt.scatter(X, y, edgecolor='r', s=20, label="Samples")
plt.xlabel("x")
plt.ylabel("y")
plt.xlim((0, 1))
plt.ylim((-2, 2))
plt.legend(loc="best")
plt.title("Degree {} \n MSE = {:.2e} (+/- {:.2e})".format(
    degrees, -scores.mean(), scores.std()))
plt.show()
```



Optimal degree in terms of MSE is:

$7.86 \times 10^{-1}$

Explain what happens when the value is higher:

*I tried the degree from 1 - 20 for each model and evaluate them. The figure is shown below. I plot the line chart for the relationships between the different degrees' model and MSE score. When the value gets bigger and bigger, I find that the MSE increases more and more significantly. Besides, according to the changes of the above figure in different degree, when the degree become higher, the model become more and more overfitting. (You can change the degree to 20 to re-run the two previous code blocks and observe the changes of figure compared with it in degree of 10)*

Explain what happens when the value is lower:

*I tried the degree from 1 - 20 for each model and evaluate them. The figure is shown below. I plot the line chart for the relationships between the different degrees' model and MSE score. When the value is lower, I find that there is no significant changes for MSE score. However, the model will become more and more underfitting according to the changes in the above figure. (You can change the degree to 1 to re-run the two previous code blocks and observe the changes of figure compared with it in degree of 10)*


(10 points)

```
MSEs = []
max_degree = 20
for degrees in range(1, max_degree+1):
    polynomial_features = PolynomialFeatures(degree=degrees,
                                              include_bias=False)

    linear_regression = LinearRegression()
    pipeline = Pipeline([("polynomial_features", polynomial_features),
                          ("linear_regression", linear_regression)])
    pipeline.fit(X[:, np.newaxis], y)
    scores = cross_val_score(pipeline, X[:, np.newaxis], y,
                             scoring="neg_mean_squared_error", cv=10)
    MSEs.append(-scores.mean())

print(MSEs)
plt.plot(range(1, max_degree+1), MSEs)
plt.show()
```

[0.8305954050139389, 0.11851458519109825, 0.29272186468432915, 0.037457715085863215, 0.03526542575950938, 0.07419619836282257, 0.6033931198803637, 1.0237951267207899, 1.0787537015790272, 0.7859477666424505, 23.62751682435972, 1731.3977790342174, 217.08211226014214, 2928.887541125662, 90537.6875987492, 49538.2595917656, 553293.286241847, 6462620.825542553, 38212884.22396408, 47978734.13649492]



## Bouns: Implement kNN (10 points)

```
import numpy as np

# Write your own implementation of KNN. Your code should pass tests in
# the cell below.
class KNN:
    def __init__(self, k=3):
        """
        Method is called when an object is created
        """
        self.k = k
        self.X_train = None
        self.y_train = None

    def fit(self, X_train, y_train):
        """
        Method is called to train (fit) the model
        """
        # Your code here
        self.X_train = X_train
        self.y_train = y_train
        print("====LabelInfo====")
```

```

print(f"X_shape: {X_train.shape}")
print(f"Y_shape: {y_train.shape}")
print("=====")
return

def predict(self, X_test) -> np.ndarray:
    """
    Method is called to predict labels for new data
    """
    # Your code here
    self.X_test = X_test
    self.distances = self.calculate_distance()
    print(f"Unsorted Distances: {self.distances}")
    self.predictions = np.argsort(self.distances, axis=1)
    print(f"Sorted Predictions: {self.predictions}")
    self.predictions = self.predictions[:, :self.k]
    print(f"Cutted Prediction: {self.predictions}")
    self.predictions = self.y_train[self.predictions]
    print(f"Labeled Predictions: {self.predictions}")
    self.predictions = np.array([np.bincount(prediction) for
prediction in self.predictions])
    print(f"Counted Predictions: {self.predictions}")
    self.predictions = np.argmax(self.predictions, axis=1)
    print(f"Final Predictions: {self.predictions}")
    return self.predictions # predictions should be in format of
Numpy array

def calculate_distance(self):
    # X_train_2 = self.X_train**2
    # X_test_2 = self.X_test**2
    # result = np.matmul(X_test_2, X_train_2.T)
    # print(X_train_2)
    # print(X_test_2)

    result = []

    for sample in self.X_test:
        # print(self.X_train - sample)
        result.append(self.X_train - sample)

    # print(result)
    result = np.array(result)**2
    # print(result)
    result = result.sum(axis=2)
    # print(result)
    return result

def basicTest(k=3):
    model = KNN(k=k)

```

```

# Basic Test
X_train = np.array([[1, 2], [2, 3], [3, 4], [5, 6], [6, 7]])
y_train = np.array([0, 0, 0, 1, 1])
X_test = np.array([[4, 5], [7, 8]])
model.fit(X_train, y_train)
predictions = model.predict(X_test)
assert np.array_equal(predictions, np.array([0, 1])), f"Expected
[0, 1] but got {predictions}"

basicTest()

=====LabelInfo=====
X_shape: (5, 2)
Y_shape: (5,)
=====
Unsorted Distances: [[18  8  2  2  8]
 [72 50 32  8  2]]
Sorted Predictions: [[2 3 1 4 0]
 [4 3 2 1 0]]
Cutted Prediction: [[2 3 1]
 [4 3 2]]
Labeled Predictions: [[0 1 0]
 [1 1 0]]
Counted Predictions: [[2 1]
 [1 2]]
Final Predictions: [0 1]

```

Run the following cell to test your code.

Note: This action is required for grading purposes (10 points)

```

def test_knn():
    model = KNN(k=3)

    # Basic Test
    X_train = np.array([[1, 2], [2, 3], [3, 4], [5, 6], [6, 7]])
    y_train = np.array([0, 0, 0, 1, 1])
    X_test = np.array([[4, 5], [7, 8]])
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    assert np.array_equal(predictions, np.array([0, 1])), f"Expected
[0, 1] but got {predictions}"

    # Edge Case Test: All points are equidistant from test point
    X_train = np.array([[1, 1], [1, 1], [1, 1]])
    y_train = np.array([0, 1, 2])
    X_test = np.array([[1, 1]])
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)
    assert predictions[0] in y_train, f"Prediction should be one of

```

the training labels, got {predictions[0]}"

*# Different k-values Test*

X\_train = np.array([[1, 2], [2, 3], [3, 4], [5, 6], [6, 7]])

y\_train = np.array([0, 0, 0, 1, 1])

X\_test = np.array([[4, 5]])

model = KNN(k=1)

model.fit(X\_train, y\_train)

predictions = model.predict(X\_test)

assert predictions[0] == 0, f"Expected 0 but got {predictions[0]}"

model = KNN(k=5)

model.fit(X\_train, y\_train)

predictions = model.predict(X\_test)

assert predictions[0] == 0, f"Expected 0 but got {predictions[0]}"

*# Tie Case Test*

X\_train = np.array([[1, 2], [2, 3], [5, 6], [6, 7]])

y\_train = np.array([0, 0, 1, 1])

X\_test = np.array([[4, 5]])

model = KNN(k=3)

model.fit(X\_train, y\_train)

predictions = model.predict(X\_test)

assert predictions[0] == 0 or predictions[0] == 1, f"Expected 0 or 1 but got {predictions[0]}"

print("All tests passed! Your code works as expected")

test\_knn()

=====LabelInfo=====

X\_shape: (5, 2)

Y\_shape: (5,)

=====

Unsorted Distances: [[18 8 2 2 8]

[72 50 32 8 2]]

Sorted Predictions: [[2 3 1 4 0]

[4 3 2 1 0]]

Cutted Prediction: [[2 3 1]

[4 3 2]]

Labeled Predictions: [[0 1 0]

[1 1 0]]

Counted Predictions: [[2 1]

[1 2]]

Final Predictions: [0 1]

=====LabelInfo=====

X\_shape: (3, 2)

Y\_shape: (3,)

=====

```

Unsorted Distances: [[0 0 0]]
Sorted Predictions: [[0 1 2]]
Cutted Prediction: [[0 1 2]]
Labeled Predictions: [[0 1 2]]
Counted Predictions: [[1 1 1]]
Final Predictions: [0]
=====LabelInfo=====
X_shape: (5, 2)
Y_shape: (5,)
=====
Unsorted Distances: [[18 8 2 2 8]]
Sorted Predictions: [[2 3 1 4 0]]
Cutted Prediction: [[2]]
Labeled Predictions: [[0]]
Counted Predictions: [[1]]
Final Predictions: [0]
=====LabelInfo=====
X_shape: (5, 2)
Y_shape: (5,)
=====
Unsorted Distances: [[18 8 2 2 8]]
Sorted Predictions: [[2 3 1 4 0]]
Cutted Prediction: [[2 3 1 4 0]]
Labeled Predictions: [[0 1 0 1 0]]
Counted Predictions: [[3 2]]
Final Predictions: [0]
=====LabelInfo=====
X_shape: (4, 2)
Y_shape: (4,)
=====
Unsorted Distances: [[18 8 2 8]]
Sorted Predictions: [[2 1 3 0]]
Cutted Prediction: [[2 1 3]]
Labeled Predictions: [[1 0 1]]
Counted Predictions: [[1 2]]
Final Predictions: [1]
All tests passed! Your code works as expected

```